

# Goal : Create a model to predict whether or not a customer will Churn .

```
In [1]: ▶ # importing basic libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: ▶ # Loading data
df = pd.read_csv(r'Telco-Customer-Churn.csv')
```

## Data Exploration

```
In [3]: ▶ df.shape
```


Out[3]: (7032, 21)

```
In [4]: ▶ df.head()
```

Out[4]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleL
0	7590-VHVEG	Female	0	Yes	No	1	No	No pl se
1	5575-GNVDE	Male	0	No	No	34	Yes	
2	3668-QPYBK	Male	0	No	No	2	Yes	
3	7795-CFOCW	Male	0	No	No	45	No	No pl se
4	9237-HQITU	Female	0	No	No	2	Yes	

5 rows × 21 columns




```
In [5]: ▶ df.tail()
```

In [5]: `df.tail()`

Out[5]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	Multiplan
7027	6840-RESVB	Male	0	Yes	Yes	24	Yes	
7028	2234-XADUH	Female	0	Yes	Yes	72	Yes	
7029	4801-JZAZL	Female	0	Yes	Yes	11	No	N
7030	8361-LTMKD	Male	1	Yes	No	4	Yes	
7031	3186-AJIEK	Male	0	No	No	66	Yes	

5 rows × 21 columns




In [6]: `df.sample(5)`

Out[6]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	Multiplan
811	3551-GA EGL	Male	0	Yes	Yes	34	No	N
4923	4298-OYIFC	Male	0	Yes	No	15	Yes	
6947	3078-ZKNTS	Female	0	Yes	Yes	13	Yes	
2664	4659-NZRUF	Female	0	No	No	19	Yes	
5183	5982-PSMKW	Female	0	Yes	Yes	23	Yes	

5 rows × 21 columns



## Quick Data Check

In [7]: ▶ df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7032 entries, 0 to 7031
Data columns (total 21 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   customerID            7032 non-null   object
 1   gender                7032 non-null   object
 2   SeniorCitizen         7032 non-null   int64
 3   Partner               7032 non-null   object
 4   Dependents            7032 non-null   object
 5   tenure               7032 non-null   int64
 6   PhoneService          7032 non-null   object
 7   MultipleLines         7032 non-null   object
 8   InternetService       7032 non-null   object
 9   OnlineSecurity        7032 non-null   object
10   OnlineBackup          7032 non-null   object
11   DeviceProtection      7032 non-null   object
12   TechSupport           7032 non-null   object
13   StreamingTV           7032 non-null   object
14   StreamingMovies       7032 non-null   object
15   Contract              7032 non-null   object
16   PaperlessBilling      7032 non-null   object
17   PaymentMethod         7032 non-null   object
18   MonthlyCharges        7032 non-null   float64
19   TotalCharges          7032 non-null   float64
20   Churn                 7032 non-null   object
dtypes: float64(2), int64(2), object(17)
memory usage: 1.1+ MB
```

## Statistical summary

In [8]: ▶ df.describe()

Out[8]:

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges
count	7032.000000	7032.000000	7032.000000	7032.000000
mean	0.162400	32.421786	64.798208	2283.300441
std	0.368844	24.545260	30.085974	2266.771362
min	0.000000	1.000000	18.250000	18.800000
25%	0.000000	9.000000	35.587500	401.450000
50%	0.000000	29.000000	70.350000	1397.475000
75%	0.000000	55.000000	89.862500	3794.737500
max	1.000000	72.000000	118.750000	8684.800000

In [9]: ▶ df.nunique()

```
In [9]: df.nunique()
```

```
Out[9]: customerID      7032
gender                2
SeniorCitizen         2
Partner               2
Dependents            2
tenure                72
PhoneService          2
MultipleLines         3
InternetService       3
OnlineSecurity        3
OnlineBackup          3
DeviceProtection      3
TechSupport           3
StreamingTV           3
StreamingMovies       3
Contract              3
PaperlessBilling      2
PaymentMethod         4
MonthlyCharges        1584
TotalCharges          6530
Churn                 2
dtype: int64
```

## Part 2: Exploratory Data Analysis

### General Feature Exploration


Confirming that there are no NaN cells by displaying NaN values per feature column.

```
In [10]: df.isna().sum()
```

```
Out[10]: customerID      0
gender                0
SeniorCitizen         0
Partner               0
Dependents            0
tenure                0
PhoneService          0
MultipleLines         0
InternetService       0
OnlineSecurity        0
OnlineBackup          0
DeviceProtection      0
TechSupport           0
StreamingTV           0
StreamingMovies       0
Contract              0
PaperlessBilling      0
PaymentMethod         0
MonthlyCharges        0
TotalCharges          0
Churn                 0
dtype: int64
```

## Checking the balance of the class labels (Churn) with a Count Plot.

```
In [11]: ▶ sns.countplot(data=df, x='Churn', hue='Churn');
```



- The churn rate is 28%. This is a significant percentage of customers, and it is important to understand why these customers are churning.

```
In [12]: ▶ df.value_counts('Churn')
```

```
Out[12]: Churn
No      5163
Yes     1869
Name: count, dtype: int64
```

```
In [13]: ▶ df.value_counts('SeniorCitizen')
```

```
Out[13]: SeniorCitizen
0      5890
1     1142
Name: count, dtype: int64
```

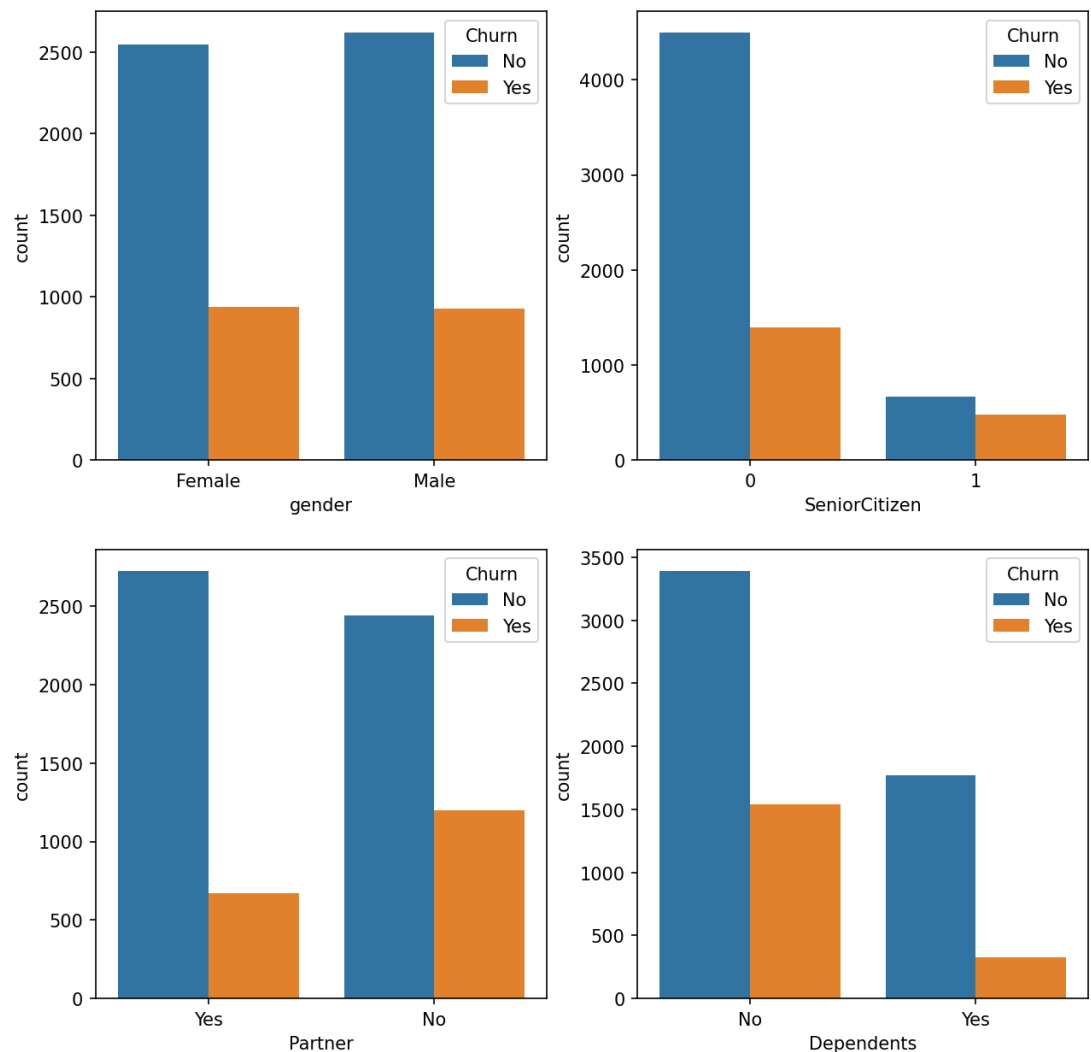
```
In [14]: ▶ fig, axes = plt.subplots(2, 2, figsize=(10, 10), dpi=150)
```

```
In [14]: fig, axes = plt.subplots(2, 2, figsize=(10, 10),dpi=150)

fig.suptitle('Countplot for Data')


sns.countplot(ax=axes[0, 0], data=df, x='gender',hue='Churn')
sns.countplot(ax=axes[0, 1], data=df, x='SeniorCitizen',hue='Churn')
sns.countplot(ax=axes[1, 0], data=df, x='Partner',hue='Churn')
sns.countplot(ax=axes[1, 1], data=df, x='Dependents',hue='Churn');
```

Countplot for Data




## Exploring the distribution of TotalCharges between Churn categories with a Box Plot or Violin Plot.\*\*

```
In [15]: ▶ sns.violinplot(data=df,x='Churn',y='TotalCharges',hue='Churn');
```



## Creating a boxplot to display the distribution of TotalCharges per Contract type,

```
In [16]: ▶ plt.figure(figsize=(7,4),dpi=150)
sns.boxplot(data=df, x='Contract',y='TotalCharges',hue='Churn')
plt.legend(loc=(1.05,0.5));
```



- On Month to Month Plan, we can say those customers are trying services but at the same time We need to analyze One-year and Two-year contract customers. We can clearly see that customers who churn are paying slightly higher charges than customers who do not


time. We need to analyze One-year and Two-year contract customers. We can clearly see that customers who churn are paying slightly higher charges than customers who do not churn. After one-year or two-year contracts, people are more likely to churn if they are having more total charge.

## Creating a bar plot to showcase the correlation of the following features

```
In [17]: ▶ corr_df = pd.get_dummies(df[['gender', 'SeniorCitizen', 'Partner', 'Depend',
'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport', 'Int
'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling', 'Paym
```

```
In [18]: ▶ corr_yes_churn = corr_df['Churn_Yes'].sort_values().iloc[1:-1]
```

```
In [19]: ▶ plt.figure(figsize=(10,4),dpi=150)
plt.title('Feature correlation with Yes Churn')
plt.xlabel('Features')
plt.ylabel('Correlation Percentage')
sns.barplot(x=corr_yes_churn.index,y=corr_yes_churn.values,hue=corr_yes_ch
plt.xticks(rotation=90);
```



## Part 3: Churn Analysis

This section focuses on segmenting customers based on their tenure, creating "cohorts", allowing us to examine differences between customer cohort segments.




## Checking available contract types

```
In [20]: df['Contract'].unique()
```

```
Out[20]: array(['Month-to-month', 'One year', 'Two year'], dtype=object)
```

**Creating a histogram for displaying the distribution of 'tenure' column, which is the amount of months a customer was or has been on a customer.**

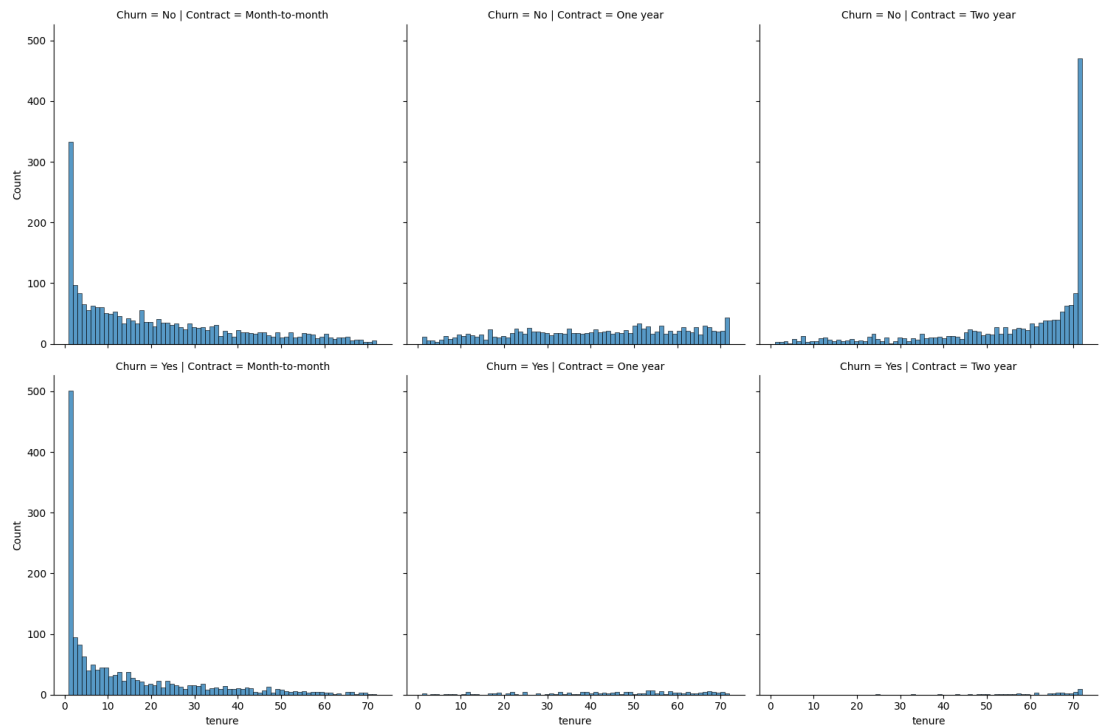
```
In [21]: plt.figure(figsize=(6,3),dpi=150)
sns.histplot(data=df,x='tenure',bins=50);
```



- Many customers have short tenures, such as one or two months.
- Many customers are short-term customers.
- There are spikes in customer tenure at regular intervals, such as 12, 24, 36, and 48 months, which suggests that these are annual plans.

```
In [22]: sns.displot(data=df,x='tenure'.bins=70,col='Contract'.row='Churn');
```

```
In [22]: sns.displot(data=df,x='tenure',bins=70,col='Contract',row='Churn');
```




- Customers with one-year and two-year contracts are not churning.
- Customers with one-month contracts are churning at a high rate.

## Creating a scatter plot of Total Charges versus Monthly Charges, and color hue by Churn.

```
In [23]: plt.figure(figsize=(10,4),dpi=150)
sns.scatterplot(data=df,x='MonthlyCharges', y='TotalCharges', hue='Churn',
```

```
Out[23]: <Axes: xlabel='MonthlyCharges', ylabel='TotalCharges'>
```



- Many customers are more likely to cancel their subscriptions when their monthly charges are higher.

## Creating Cohorts based on Tenure

Let's begin by treating each unique tenure length, 1 month, 2 month, 3 month...N months as its own cohort.

### Calculating the Churn rate (percentage that had Yes Churn) per cohort.

```
In [24]: ► yes_churn = df.groupby(['Churn', 'tenure']).count().transpose()['Yes']
```

```
In [25]: ► yes_churn
```

Out[25]:

	tenure	1	2	3	4	5	6	7	8	9	10	...	63	64	65	66	67	68	69
Churn Data	customerID	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8
	gender	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8
	SeniorCitizen	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8
	Partner	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8
	Dependents	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8
	PhoneService	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8
	MultipleLines	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8
	InternetService	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8
	OnlineSecurity	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8
	OnlineBackup	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8
	DeviceProtection	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8
	TechSupport	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8
	StreamingTV	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8
	StreamingMovies	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8
	Contract	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8
	PaperlessBilling	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8
	PaymentMethod	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8
	MonthlyCharges	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8
	TotalCharges	380	123	94	83	64	40	51	42	46	45	...	4	4	9	13	10	9	8

19 rows × 72 columns



```
In [26]: ► no_churn = df.groupby(['Churn', 'tenure']).count().transpose()['No']
```

```
In [27]: ► no_churn
```

In [27]: ▶ no\_churn

Out[27]:

	tenure	1	2	3	4	5	6	7	8	9	10	...	63	64	65	66	67	68	6
CustomerID	customerID	233	115	106	93	69	70	80	81	73	71	...	68	76	67	76	88	91	8
	gender	233	115	106	93	69	70	80	81	73	71	...	68	76	67	76	88	91	8
	SeniorCitizen	233	115	106	93	69	70	80	81	73	71	...	68	76	67	76	88	91	8
	Partner	233	115	106	93	69	70	80	81	73	71	...	68	76	67	76	88	91	8
	Dependents	233	115	106	93	69	70	80	81	73	71	...	68	76	67	76	88	91	8
	PhoneService	233	115	106	93	69	70	80	81	73	71	...	68	76	67	76	88	91	8
	MultipleLines	233	115	106	93	69	70	80	81	73	71	...	68	76	67	76	88	91	8
	InternetService	233	115	106	93	69	70	80	81	73	71	...	68	76	67	76	88	91	8
	OnlineSecurity	233	115	106	93	69	70	80	81	73	71	...	68	76	67	76	88	91	8
	OnlineBackup	233	115	106	93	69	70	80	81	73	71	...	68	76	67	76	88	91	8
	DeviceProtection	233	115	106	93	69	70	80	81	73	71	...	68	76	67	76	88	91	8
	TechSupport	233	115	106	93	69	70	80	81	73	71	...	68	76	67	76	88	91	8
	StreamingTV	233	115	106	93	69	70	80	81	73	71	...	68	76	67	76	88	91	8
	StreamingMovies	233	115	106	93	69	70	80	81	73	71	...	68	76	67	76	88	91	8
	Contract	233	115	106	93	69	70	80	81	73	71	...	68	76	67	76	88	91	8
	PaperlessBilling	233	115	106	93	69	70	80	81	73	71	...	68	76	67	76	88	91	8
	PaymentMethod	233	115	106	93	69	70	80	81	73	71	...	68	76	67	76	88	91	8
	MonthlyCharges	233	115	106	93	69	70	80	81	73	71	...	68	76	67	76	88	91	8
	TotalCharges	233	115	106	93	69	70	80	81	73	71	...	68	76	67	76	88	91	8

19 rows × 72 columns



In [28]: ▶ churn\_rate = 100 \* yes\_churn / (no\_churn + yes\_churn)


In [29]: ▶ churn\_rate.transpose()

```
In [29]: churn_rate.transpose()
```


Out[29]:

	customerID	gender	SeniorCitizen	Partner	Dependents	PhoneService	Multipl
tenure							
1	61.990212	61.990212	61.990212	61.990212	61.990212	61.990212	61.9
2	51.680672	51.680672	51.680672	51.680672	51.680672	51.680672	51.6
3	47.000000	47.000000	47.000000	47.000000	47.000000	47.000000	47.0
4	47.159091	47.159091	47.159091	47.159091	47.159091	47.159091	47.1
5	48.120301	48.120301	48.120301	48.120301	48.120301	48.120301	48.1
...	...	...	...	...	...	...	...
68	9.000000	9.000000	9.000000	9.000000	9.000000	9.000000	9.0
69	8.421053	8.421053	8.421053	8.421053	8.421053	8.421053	8.4
70	9.243697	9.243697	9.243697	9.243697	9.243697	9.243697	9.2
71	3.529412	3.529412	3.529412	3.529412	3.529412	3.529412	3.5
72	1.657459	1.657459	1.657459	1.657459	1.657459	1.657459	1.6

72 rows × 19 columns



```
In [30]: churn_rate.transpose()['customerID'].plot();
```



- Customer churn rate decreases as tenure increases.

## Broader Cohort Groups

Based on the tenure column values, creating a new column called Tenure Cohort that creates 4 separate categories:

- '0-12 Months'
- '12-24 Months'
- '24-48 Months'
- 'Over 48 Months'

```
In [31]: ► def cohort(tenure):  
          if tenure<13:  
              return '0-12 Months'  
          elif tenure < 25:  
              return '12-24 Months'  
          elif tenure<49:  
              return '24-48 Months'  
          else:  
              return 'Over 48 Months'
```

```
In [32]: ► df['Tenure Cohort'] = df['tenure'].apply(cohort)
```

```
In [33]: ► df[['tenure', 'Tenure Cohort']]
```


Out[33]:

	tenure	Tenure Cohort
0	1	0-12 Months
1	34	24-48 Months
2	2	0-12 Months
3	45	24-48 Months
4	2	0-12 Months
...	...	...
7027	24	12-24 Months
7028	72	Over 48 Months
7029	11	0-12 Months
7030	4	0-12 Months
7031	66	Over 48 Months

7032 rows × 2 columns

## Scatterplot of Total Charges versus Monthly Charts,

```
In [34]: ▶ plt.figure(figsize=(10,4),dpi=200)
sns.scatterplot(data=df,x='MonthlyCharges',y='TotalCharges',hue='Tenure Co
```

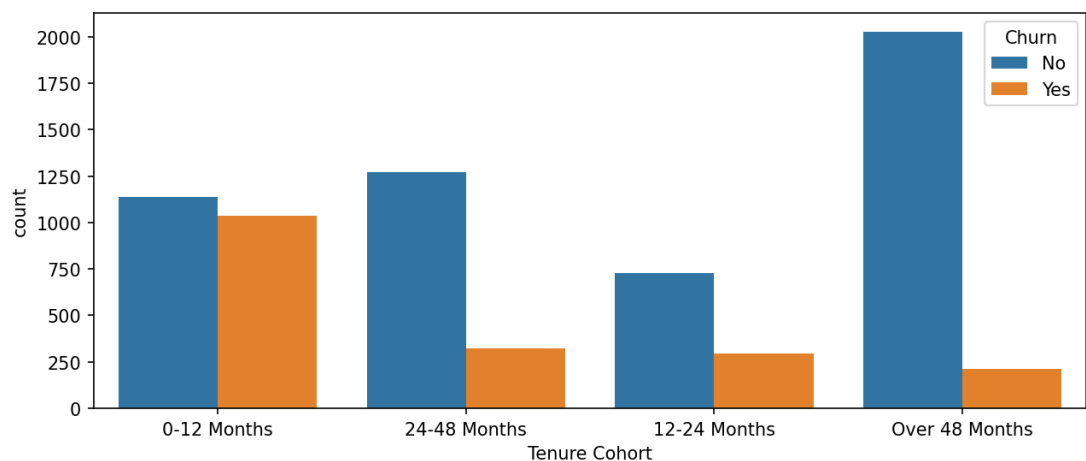


- Cohorts with longer tenures have the highest total charges.
- We can also see that there is a drop in total charges for some cohorts immediately after their monthly charges increase.

## Count plot showing the churn count per cohort.

```
In [35]: ▶ plt.figure(figsize=(10,4),dpi=150)
sns.countplot(data=df,x='Tenure Cohort',hue='Churn')
```

Out[35]: <Axes: xlabel='Tenure Cohort', ylabel='count'>




- Customers who stay for more than 48 months are less likely to churn.
- Customers with only 0-48 months of service are most likely to churn.

Create a grid of Count Plots showing counts per Tenure Cohort, separated out by contract type and colored by the Churn hue.

```
In [36]: ▶ sns.catplot(data=df, x='Tenure Cohort', col='Contract', hue='Churn', kind='count')
```

```
In [36]: sns.catplot(data=df, x='Tenure Cohort', col='Contract', hue='Churn',kind='')
```



## Part 4: Predictive Modeling

\*\*Let's explore 4 different tree based methods: A Single Decision Tree, Random Forest, AdaBoost, Gradient Boosting.

### Data Splitting

```
In [37]: X = df.drop(['Churn', 'customerID'],axis=1)
X = pd.get_dummies(X,dtype=int,drop_first=True)

y = df['Churn']
```

```
In [38]: from sklearn.model_selection import train_test_split
```

```
In [39]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1, r
```

```
In [40]: print("Shape of X_train: ",X_train.shape)
print("Shape of X_test: ", X_test.shape)
print("Shape of y_train: ",y_train.shape)
print("Shape of y_test:",y_test.shape)
```

```
Shape of X_train: (6328, 33)
Shape of X_test: (704, 33)
Shape of y_train: (6328,)
Shape of y_test: (704,)
```

### Model Selection : Single Decision Tree

```
In [41]: from sklearn.tree import DecisionTreeClassifier
```

```
In [42]: dt = DecisionTreeClassifier(max_depth=6)
```

```
In [43]: dt.fit(X_train, y_train)
```



```
In [43]: dt.fit(X_train, y_train)
```

```
Out[43]: DecisionTreeClassifier  
DecisionTreeClassifier(max_depth=6)
```

## Model Evaluation

```
In [44]: from sklearn.metrics import accuracy_score, confusion_matrix, ConfusionMat
```

```
In [45]: preds = dt.predict(X_test)
```


```
In [46]: accuracy_score(y_test, preds)
```

```
Out[46]: 0.8082386363636364
```

```
In [47]: confusion_matrix(y_test, preds)
```

```
Out[47]: array([[497,  60],  
               [ 75,  72]], dtype=int64)
```

```
In [48]: ConfusionMatrixDisplay(confusion_matrix(y_test, preds), display_labels=dt.c
```




```
In [49]: print(classification_report(y_test, preds))
```

In [49]: `print(classification_report(y_test,preds))`

	precision	recall	f1-score	support
No	0.87	0.89	0.88	557
Yes	0.55	0.49	0.52	147
accuracy			0.81	704
macro avg	0.71	0.69	0.70	704
weighted avg	0.80	0.81	0.80	704

- The overall accuracy of the tree model is 81%, which is good.
- The precision for the "No" class is 87%, which means that 87% of the customers that the model predicted would not churn actually did not churn.
- The recall for the "No" class is 89%, which means that the model identified 89% of all customers who did not churn.
- The precision for the "Yes" class is 55%, which means that 55% of the customers that the model predicted would churn actually did churn.
- The recall for the "Yes" class is 49%, which means that the model identified 49% of all customers who did churn.

In [50]: `imp_feats = pd.DataFrame(data=dt.feature_importances_,index=X.columns,columns=)`



In [51]: `imp_feats`

```
In [51]: ▶ imp_feats
```

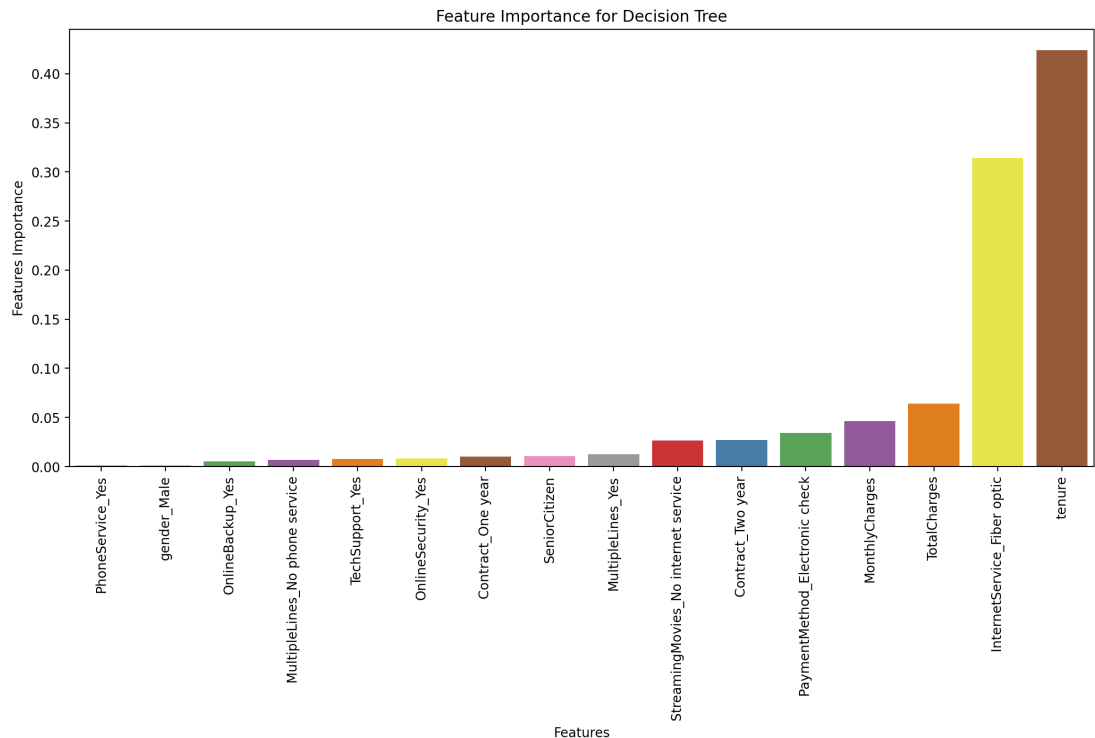
```
Out[51]:
```

	Feature Importance
DeviceProtection_No internet service	0.000000
Tenure Cohort_12-24 Months	0.000000
PaymentMethod_Mailed check	0.000000
PaymentMethod_Credit card (automatic)	0.000000
PaperlessBilling_Yes	0.000000
StreamingMovies_Yes	0.000000
StreamingTV_Yes	0.000000
StreamingTV_No internet service	0.000000
TechSupport_No internet service	0.000000
DeviceProtection_Yes	0.000000
Tenure Cohort_24-48 Months	0.000000
OnlineBackup_No internet service	0.000000
Tenure Cohort_Over 48 Months	0.000000
InternetService_No	0.000000
Dependents_Yes	0.000000
OnlineSecurity_No internet service	0.000000
Partner_Yes	0.000000
PhoneService_Yes	0.000890
gender_Male	0.001237
OnlineBackup_Yes	0.005341
MultipleLines_No phone service	0.006962
TechSupport_Yes	0.007868
OnlineSecurity_Yes	0.008376
Contract_One year	0.010021
SeniorCitizen	0.010825
MultipleLines_Yes	0.012432
StreamingMovies_No internet service	0.026290
Contract_Two year	0.027065
PaymentMethod_Electronic check	0.034436
MonthlyCharges	0.046115
TotalCharges	0.064168
InternetService_Fiber optic	0.314060
tenure	0.423914

```
In [52]: ▶ imp_feats = imp_feats[imp_feats['Feature Importance'] > 0]
```

```
In [53]: ▶ plt.figure(figsize=(14,6)).dpi=200)
```

```
In [53]: ▶ plt.figure(figsize=(14,6),dpi=200)
sns.barplot(data=imp_feats,x=imp_feats.index,y='Feature Importance',hue='F
plt.title('Feature Importance for Decision Tree')
plt.legend().remove()
plt.xlabel('Features')
plt.ylabel('Features Importance')
plt.xticks(rotation=90);
```

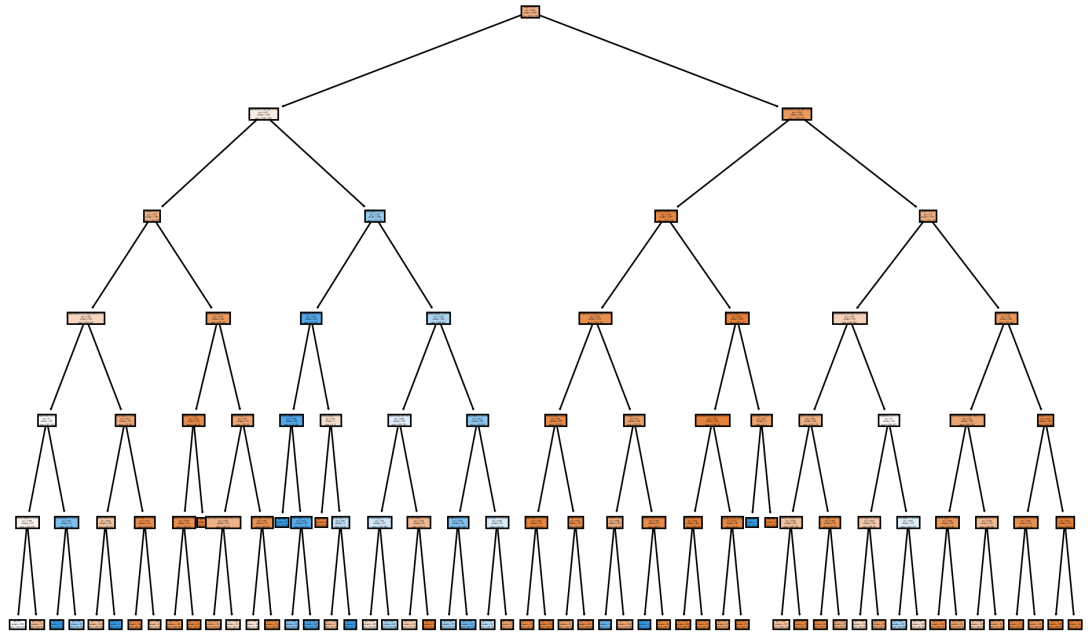


- Tenure is the most important feature for predicting customer churn. Customers who have been with the company for a longer period of time are less likely to churn.
- Customers with fiber optic internet service are less likely to churn than customers with other types of internet service. This suggests that fiber optic internet service is a valuable service that customers are willing to pay for.
- Customers with higher total and monthly charges are more likely to churn. This suggests that customers are price-sensitive and are looking for the best possible value.
- Customers who pay their bills with electronic checks are more likely to churn than customers who pay with other methods. This suggests that electronic check payments may be associated with financial problems, which can lead to churn.
- Customers who need tech support and customers who have contracts that are two years long are more likely to churn. This suggests that these customers may be having problems with the company's products or services.

```
In [54]: ▶ from sklearn.tree import plot_tree
```

```
In [54]: ▶ from sklearn.tree import plot_tree

plt.figure(figsize=(12,8),dpi=200)
plot_tree(dt,filled=True,feature_names=X.columns);
```



## Model Selection : Random Forest

```
In [55]: ▶ from sklearn.ensemble import RandomForestClassifier
```

```
In [56]: ▶ rf = RandomForestClassifier()
```

```
In [57]: ▶ rf.fit(X_train,y_train)
```

```
Out[57]: ▼ RandomForestClassifier
RandomForestClassifier()
```

## Model Evaluation

```
In [58]: ▶ preds = rf.predict(X_test)
```

```
In [59]: ▶ accuracy_score(y_test,preds)
```


```
Out[59]: 0.7869318181818182
```

```
In [60]: ▶ confusion_matrix(y_test,preds)
```

```
Out[60]: array([[490,  67],
                [ 83,  64]], dtype=int64)
```

```
In [61]: ▶ ConfusionMatrixDisplay(confusion_matrix(y_test,preds).display_labels=rf.c
```

In [61]: `ConfusionMatrixDisplay(confusion_matrix(y_test,preds), display_labels=rf.c`



In [62]: `print(classification_report(y_test,preds))`

	precision	recall	f1-score	support
No	0.86	0.88	0.87	557
Yes	0.49	0.44	0.46	147
accuracy			0.79	704
macro avg	0.67	0.66	0.66	704
weighted avg	0.78	0.79	0.78	704

- The model is performing well overall, with an accuracy of 81%.
- The model is performing better on the "No" class than on the "Yes" class.
- This could be due to a number of factors, such as imbalanced data, insufficient features, or overfitting.
- To improve the performance of the model on the "Yes" class, you can try balancing the dataset, adding more features, using a different machine learning algorithm, or understanding which features are most important for predicting customer churn.

**Tune the max depth parameter of the model to see if it improves accuracy.**

In [63]: `rf = RandomForestClassifier(max_depth=6)`

In [64]: `rf.fit(X_train,y_train)`

```
In [64]: rf.fit(X_train,y_train)
```

```
Out[64]: RandomForestClassifier  
RandomForestClassifier(max_depth=6)
```

## Model Evaluation

```
In [65]: preds = rf.predict(X_test)
```


```
In [66]: accuracy_score(y_test,preds)
```

```
Out[66]: 0.8210227272727273
```

```
In [67]: confusion_matrix(y_test,preds)
```

```
Out[67]: array([[512,  45],  
               [ 81,  66]], dtype=int64)
```

```
In [68]: ConfusionMatrixDisplay(confusion_matrix(y_test,preds), display_labels=rf.c
```



```
In [69]: print(classification_report(y_test,preds))
```

```
In [69]: print(classification_report(y_test,preds))
```

	precision	recall	f1-score	support
No	0.86	0.92	0.89	557
Yes	0.59	0.45	0.51	147
accuracy			0.82	704
macro avg	0.73	0.68	0.70	704
weighted avg	0.81	0.82	0.81	704

- Overall, the RandomForestClassifier(max\_depth=6) model is performing well, with an accuracy of 83%. It is also performing better on the "Yes" class than the previous model with a max\_depth of 3. This suggests that increasing the max\_depth parameter has helped to improve the model's ability to identify customers who are at risk of churning.
- However, it is important to note that the model is still making more false positives than false negatives. This means that the model is more likely to incorrectly predict that a customer will churn than to incorrectly predict that a customer will not churn.

## Model Selection : Gradient Boosting

```
In [70]: from sklearn.ensemble import GradientBoostingClassifier
```

```
In [71]: gb_model = GradientBoostingClassifier()
```

```
In [72]: gb_model.fit(X_train,y_train)
```

```
Out[72]: ▾ GradientBoostingClassifier
          GradientBoostingClassifier()
```

## Model Evaluation

```
In [73]: gb_preds = gb_model.predict(X_test)
```

```
In [74]: accuracy_score(y_test,gb_preds)
```

```
Out[74]: 0.8181818181818182
```


```
In [75]: confusion_matrix(y_test,gb_preds)
```

```
Out[75]: array([[503,  54],
                [ 74,  73]], dtype=int64)
```

```
In [76]: ConfusionMatrixDisplay(confusion_matrix(y_test,gb_preds).display_labels=
```



In [76]: `ConfusionMatrixDisplay(confusion_matrix(y_test,gb_preds), display_labels=g`



In [77]: `print(classification_report(y_test,gb_preds))`

	precision	recall	f1-score	support
No	0.87	0.90	0.89	557
Yes	0.57	0.50	0.53	147
accuracy			0.82	704
macro avg	0.72	0.70	0.71	704
weighted avg	0.81	0.82	0.81	704

- Overall, the GradientBoostingClassifier model is performing well, with an accuracy of 82%. It is also performing better on the "Yes" class than the RandomForestClassifier model, with a precision of 0.57 and a recall of 0.50. This suggests that the GradientBoostingClassifier model is better able to identify customers who are at risk of churning.
- However, it is important to note that the model is still making more false positives than false negatives. This means that the model is more likely to incorrectly predict that a customer will churn than to incorrectly predict that a customer will not churn.

## Model Selection : Adaptive Boosting

In [78]: `from sklearn.ensemble import AdaBoostClassifier`

In [79]: `ada_model = AdaBoostClassifier()`

```
In [79]: ▶ ada_model = AdaBoostClassifier()
```

```
In [80]: ▶ ada_model.fit(X_train,y_train)
```

```
Out[80]: ▼ AdaBoostClassifier  
AdaBoostClassifier()
```

## Model Evaluation

```
In [81]: ▶ ada_preds = ada_model.predict(X_test)
```


```
In [82]: ▶ accuracy_score(y_test,ada_preds)
```

```
Out[82]: 0.8295454545454546
```

```
In [83]: ▶ confusion_matrix(y_test,ada_preds)
```

```
Out[83]: array([[504,  53],  
               [ 67,  80]], dtype=int64)
```

```
In [84]: ▶ ConfusionMatrixDisplay(confusion_matrix(y_test,ada_preds), display_labels=
```



```
In [85]: ▶ print(classification_report(y_test,ada_preds))
```

```
In [85]: ▶ print(classification_report(y_test,ada_preds))
```

	precision	recall	f1-score	support
No	0.88	0.90	0.89	557
Yes	0.60	0.54	0.57	147
accuracy			0.83	704
macro avg	0.74	0.72	0.73	704
weighted avg	0.82	0.83	0.83	704

- Overall, the AdaBoostClassifier model is performing well, with an accuracy of 83%. It is also performing better on the "Yes" class than the previous models, with a precision of 0.60 and a recall of 0.54. This suggests that the AdaBoostClassifier model is better able to identify customers who are at risk of churning.
- However, it is important to note that the model is still making more false positives than false negatives. This means that the model is more likely to incorrectly predict that a customer will churn than to incorrectly predict that a customer will not churn.

**Tune the parameter of the adaboost to see if it improves accuracy.**

```
In [86]: ▶ len(X.columns)
```

```
Out[86]: 33
```

```
In [87]: ▶ error_rates = []


for n in range(1,34):
    model = AdaBoostClassifier(n_estimators=n)
    model.fit(X_train,y_train)
    preds = model.predict(X_test)

    err = 1 - accuracy_score(y_test, preds)

    error_rates.append(err)
```

```
In [88]: ▶ plt.plot(range(1,34),error_rates)
```

```
In [88]: ▶ plt.plot(range(1,34),error_rates)
plt.grid()
```



```
In [89]: ▶ ada_model = AdaBoostClassifier(n_estimators=11)
```

```
In [90]: ▶ ada_model.fit(X_train,y_train)
```

```
Out[90]: ▼ AdaBoostClassifier
AdaBoostClassifier(n_estimators=11)
```

## Model Evaluation

```
In [91]: ▶ ada_preds = ada_model.predict(X_test)
```

```
In [92]: ▶ accuracy_score(y_test,ada_preds)
```


```
Out[92]: 0.8352272727272727
```

```
In [93]: ▶ confusion_matrix(y_test,ada_preds)
```

```
Out[93]: array([[522, 35],
               [ 81, 66]], dtype=int64)
```

```
In [94]: ▶ ConfusionMatrixDisplay(confusion_matrix(y_test,ada_preds).display_labels=
```

In [94]: `ConfusionMatrixDisplay(confusion_matrix(y_test,ada_preds), display_labels=`



In [95]: `print(classification_report(y_test,ada_preds))`

	precision	recall	f1-score	support
No	0.87	0.94	0.90	557
Yes	0.65	0.45	0.53	147
accuracy	0.84			704
macro avg	0.76	0.69	0.72	704
weighted avg	0.82	0.84	0.82	704

- Overall, the AdaBoostClassifier(n\_estimators=11) model is performing well, with an accuracy of 84%. It is also performing better on the "Yes" class than the previous models, with a precision of 0.65 and a recall of 0.45. This suggests that increasing the number of base estimators in the AdaBoost model has helped to improve the model's ability to identify customers who are at risk of churning.
- However, it is important to note that the model is still making more false positives than false negatives. This means that the model is more likely to incorrectly predict that a customer will churn than to incorrectly predict that a customer will not churn.

## Model Deployment

In [96]: `# from joblib import dump,load`

In [97]: `# dump(final_model, 'ada_model.joblib')`

```
In [97]: ► # dump(final_model, 'ada_model.joblib')
```